

Multi-Threading

- Parallelizing Independent Tasks
 - Synchronisation with Mutexes
 - One-Time Execution
 - Messaging with Condition Variables
 - Atomic Operations
 - Direct Use of Threads
 - Native Threading Model Handles
 - Concurrency Recommendations
-

With C++11 support for multi-threading was introduced.*



For more information on concurrency support in C++11 see:
<http://en.cppreference.com/w/cpp/atomic>, and
http://en.cppreference.com/w/cpp/language/memory_model

*: At first glance concurrency features may appear "as just some more library classes and functions". But beyond the hood, and especially in the area of allowed optimisations and to provide [Cache Coherence](#) on modern multi-core CPUs, concurrency is closely intertwined with code generation issues.

Parallelizing Independent Tasks

For complex tasks that can be split into independent parts, concurrency **and scalability to multiple cores** can be easily achieved by following a simple recipe:

1. Separate the task into a number of different functions (or calls of the same function with different arguments).
2. Run each such function by handing it over to `std::async`, storing the future that is returned (easiest in an auto-typed variable).
3. Fetch (and combine) the results by calling the member function `get` for each future.

That way all the functions may run concurrently and the last step synchronizes by waiting for completion.



For more information on parallelizing independent task that way see: <http://en.cppreference.com/w/cpp/thread/async>

*: Technically the return value of `std::async` is an `std::future` but as this is a template and the type is usually somewhat different to spell out, most usages of `std::async` store the result in an auto-typed variable.

Foundation: Futures and Promises

The foundation on which parallelizing tasks is build are **Futures and Promises**.

These need not be fully understood to apply the API (exemplified on the next pages), but may help to understand the basic machinery:

- A future is the concrete handle which a client can use to fetch the result, presumably made available by a different thread of execution.
- A promise is a helper class which may be used in a separate thread to make a result available for a client.



For more information on promises and futures see:
<http://en.cppreference.com/w/cpp/thread/future>
<http://en.cppreference.com/w/cpp/thread/promise>
http://en.cppreference.com/w/cpp/thread/packaged_task

Parallelizing Example

The following example calculates the sum of the first N elements in data by splitting the work of `std::accumulate`, into two separate function calls, that may run concurrently:*

```
// calculate sum of first N values in data
//
long long sum(const int data[], std::size_t N) {
    auto lower_part_sum = std::async(
        [=]{ return std::accumulate(&data[0], &data[N/2], 0LL); }
    );
    auto upper_part_sum = std::async(
        [=]{ return std::accumulate(&data[N/2], &data[N], 0LL); }
    );
    return lower_part_sum.get()
        + upper_part_sum.get();
}
```

*: Note the use of lambdas above – the actual call to `std::accumulate` will only happen when the lambda gets executed! A similar effect can be achieved as follows:

```
// preparing the callable for std::async with std::bind:
... std::async(std::bind(std::accumulate, &data[0], &data[N/2], 0LL)) ...
... std::async(std::bind(std::accumulate, &data[N/2], &data[N], 0LL)) ...
```

Default Launch Policy

The default behavior is that the system decides on its own whether an asynchronously started task is run concurrently.*

Using `std::async` **without an explicit launch policy**

- **is just a hint** that it is acceptable to run a callable unit of code concurrently,
- as long as it has finished (and possibly returned a result) latest when the `get`-call returns, which has been invoked on the `std::future`.



Be sure **not** to use the default launch policy but specify concurrent execution explicitly (see next page) whenever it is essential that two callables run concurrently.

*: It is a well known effect that too many parallel threads of execution may rather degrade performance. Especially if threads are CPU bound, it makes little sense to have more threads as cores. Therefore the standard gives considerable freedom to the implementation, which might implement concurrency with `std::async` with a thread pool and turn to synchronous execution or lazy evaluation at a certain threshold.

Explicit Launch Policies

There is a second version of `std::async` which has a first argument to specify the launch strategy.

The standard defines two values:

- `std::launch::async`
if **not set** the callable will **not** run on its own thread;^{*}
- `std::launch::deferred` if **set** the callable will **not** be called before `get` is invoked.



For more information on launch policies see:
<http://en.cppreference.com/w/cpp/thread/launch>

^{*}: One case in which this setting this flag may make sense is to test the program logic independent from possible problems created through race conditions or deadlocks, originating from dependencies between the "independent tasks", that had been overlooked.

Catching Exceptions

If the callable started via `std::async` throws an exception, it will appear as if it were thrown from the call to `get`.

Hence, if the asynchronously run task may throw, fetching the result should be done in a try block:

```
auto task1 = std::async( ... ); // whatever-is-to-do (and may throw)
auto task2 = std::async( ... ); // whatelse-is-to-do (and may throw)
...
try { ... task1.get() ... }
} catch ( ... ) { // what may be thrown from whatever-is-to-do
    ... // handle the case that whatever-is-to-do threw
}
try { ... task2.get() ... }
} catch ( ... ) { // what may be thrown from whatelse-is-to-do
    ... // handle the case that whatelse-is-to-do threw
}
```

Communication between Independent Tasks

First of all: If the need arises to communicate between independent tasks, this should be taken **as a strong warning** that such tasks are actually not independent.



If parallel tasks are not independent, further needs follow quickly with respect to synchronize access to shared data ...
with all the further intricacies following from this.*

Nevertheless there is one common case that requires a simple form of communication between otherwise independent tasks.

- If there are several tasks working towards a common goal
- of which one fails, making the goal unattainable,
- the others should not waste CPU-time needlessly.

*: In other words: Pandora's proverbial can of worms opens quickly and widely ...

Communicate Failure between Concurrent Tasks

A basic design that communicates failure between partners working towards a common goal is outlined in the following example.*

The workers could look about so ...

... being run by that code:

```
... f( ... ,
    std::atomic_bool &die) {
    for ( ... ) {
        if (die) return ...;
        ...
        ... // some complex
        ... // algorithm
        ...
        // may fail here
        if ( ... ) {
            die = true;
            return ...;
        }
    }
}
```

```
std::atomic_bool die{false};
auto task1 = std::async(
    [&die]{ f(... , die); }
);
auto task2 = std::async(
    [&die]{ f(... , die); }
);
...
... task1.get() ...
... task2.get() ...
...
if (die) {
    // goal not reached
    ...
}
```

*: To keep this code simple f just returns in case of problems, though it requires only a few changes if problems should be communicated to the caller via exceptions.

Synchronisation with Mutexes

The word *Mutex* abbreviates *Mutual Exclusion* and describe the basic purpose of the feature.

- Allow only one thread to enter a **Critical Section**, typically non-atomically executed sequence of statements
- which temporarily invalidate an **Class Invariant**, or
- in other ways accesses a resource not designed for shared use.

In general, mutexes have at least two operations* for

- **lock**-ing and
- **unlock**-ing,

but frequently provide additional features to make their practical use more convenient or less error prone.

*: Though the operations may not be spelled exactly *lock* and *unlock* ... (especially as mutexes are somewhat related to [Semaphores], which originally named their lock (-like) acquire operation *P* and their unlock (-like) release operation *V*).

Mutex Example (1)

The following example calculates one table from another one:*

```
template<typename In, typename Out, typename Transformation>
void worker(const In data[], std::size_t data_size,
            Out result[], std::atomic_size_t &total_progress,
            Transformation func) {
    static std::mutex critical_section;
    while (total_progress < data_size) {
        critical_section.lock();
        constexpr auto chunks = std::size_t{100};
        const auto beg = total_progress;
        const auto end = ((data_size - total_progress) > chunks)
            ? total_progress += chunks
            : total_progress = data_size;
        critical_section.unlock();
        std::transform(&data[beg], &data[end], &result[beg], func);
    }
}
```

*: The work is shared by any number of worker task run concurrently, each fetches and transforms a fixed number of values. This can be advantageous to splitting the work by calculating fixed-size regions of the table in advance, if the transformation function has a largely varying runtime depending on the argument value.

Mutex Example (2)

Assuming the transformation is to calculate square roots and there are two arrays of size N, say

- data (filled with values to transform), and
- sqrts to store the results

workers may be created (to be handed over to `std::async`) as follows:*

```
std::atomic_size_t processed_count{0};
auto worker_task =
    [&]() { worker(data, N, sqrts, processed_count,
                 [] (double e) { return std::sqrt(e); });
    };
```

*: Given the above, a particular nifty way to create and run workers were:

```
// assuming NCORES holds the number of cores to use by workers:
std::array<std::future<void>, NCORES> workers;
for (auto &w : workers)
    w = std::async(worker_task);
for (auto &w : workers)
    try { w.get(); } catch (...) {}
```

Mutexes and RAII

As the mutex operations **lock** and **unlock** need to come correctly paired, they make a good candidate to apply a technique called **RAII**.*

It works by creating a wrapper class, executing

- the acquiring operation (or *lock*-ing in this case) in its constructor, and
- the releasing operation (or *unlock*-ing) in its destructor.

Such helper classes are available in C++11 `std::lock_guard`.

The big advantage is that unlocking the mutex is guaranteed for code blocks defining a RAII-style (guard) object locally, no matter whether control flow reached its end, or by `break`, `return`, or some exception.



For more information on the RAII-style use of mutexes see:
http://en.cppreference.com/w/cpp/thread/lock_guard

*: This **TLA** is an abbreviation Bjarne Stroustrup once coined for *Resource Acquisition is Initialisation*. In a recent interview Stroustrup revealed that he is not particularly happy with the term he once chose. But to change it would require to travel back in a time machine and suggest something more appropriate to him, as today the term RAII is in much too widespread use to be replaced by something else.

Mutex Variants

Mutexes in C++11 come in a number of flavours, which controlling their behaviour with respect to the following details:

- Whether or not some thread that already locked a mutex may lock it once more (and needs to release it as often).
- Whether or not a thread waits if it finds a mutex locked by some other thread, and in the latter case until *when* (clock-based time point) or or *how long* (duration) it waits.

For all mutex variants there are also variants of RAII-style lock guards.



For more information on the different variants of mutexes and RAII-style wrappers see:

http://en.cppreference.com/w/cpp/thread/recursive_mutex

http://en.cppreference.com/w/cpp/thread/timed_mutex

http://en.cppreference.com/w/cpp/thread/recursive_timed_mutex and

http://en.cppreference.com/w/cpp/thread/unique_lock

C++14: Upgradable Locks

C++14 added the class `std::shared_lock`, supporting a frequent necessity:*

Multiple Reader/Single Write Locking Schemes

- Any number of (reader) threads may successfully `shared_lock` that kind of mutex ...
- ... but only one single (writer) thread is allowed to actually lock it (unshared).

C++14 also provides a RAIL-style wrapper for shared locking.



For more information on shared locking see:

http://en.cppreference.com/w/cpp/thread/shared_timed_mutex
and http://en.cppreference.com/w/cpp/thread/shared_lock

*: Note that the terms "reader" and "writer" indicates the typical use of that kind of mutex, assuming that it is sufficient for readers to obtain the lock shared, as it guarantees the invariants hold but no modifications are made, while writers will need to temporarily break invariants.

Defeating Deadlocks Caused by Mutex-Locking

As a potentially blocking mechanism mutexes are famous for creating deadlocks, i.e

- in the situation where two resources *A* and *B* are required,
- one thread acquires these in the order *first A, then B* and
- another thread acquires these in the order *first B, then A*.*

The obvious counter measure is to acquire locks always in the same order, as achievable with `std::lock` and `std::try_lock`.



For more information on locking several mutexes semantically atomic (i.e. without creating the potential for dead-locks) see: <http://en.cppreference.com/w/cpp/thread/lock> and http://en.cppreference.com/w/cpp/thread/try_lock

*: In practice, the potential for deadlocks is often not as obvious as in this example but much more intricate and close to impossible to spot, even in scrutinising code reviews. So, if (accidental) deadlocks cannot be avoided, sometimes a "self-healing" strategy is applied that works follows:

If more than one lock needs to be acquired, acquire at least all others with setting a time-out. If that hits, **release all locks** acquired so far, **delay** for some small amount of time (usually determined in a window with some slight randomness), then **try again** (and maybe in a different order).

One-Time Execution

For a particular scenario that would otherwise require the use of mutex-es to avoid a **Race Condition**, there is pre-built solution in C++11.

Executing a piece of code exactly once can be achieved in a cookbook-style as follows:

```
// in a scope reachable from all usage points:  
std::once_flag this_code_once;  
...  
std::call_once(this_code_once, ...some callable... ); // somewhere  
...  
std::call_once(this_code_once, ... ); // maybe somewhere else
```

For any of the callables associated with the same instance of an `std::once_flag` via `std::call_once`, without further protection via mutexes it is guaranteed that **at most one** is executed **at most once**.



For more information on guaranteed one time execution see:
http://en.cppreference.com/w/cpp/thread/once_flag and
http://en.cppreference.com/w/cpp/thread/call_once

One-Time Execution Example

A typical use case for guaranteed one-time execution is some initialisation, that may be expensive and is therefore delayed until a function that depends on it is called the first time.

The following fragment avoids parallel initialisations of table:

```
... foo( ... ) {  
    static std::once_flag init;  
    static std::array<int, 1000> table;  
    std::call_once(init, [&table]() {  
        ... // pre-calculate table when foo runs for the first time  
    });  
    ... //  
}
```



For **Static Local Variables** C++11 guarantees correct behaviour even when `std::call_once` is not used, as outlined here:
http://en.cppreference.com/w/cpp/language/storage_duration

Local static Initialisation

Since C++11 supports multi-threading in the core language, initialising local static variables is protected to be executed at most once:

```
... foo( ... ) {  
    static const int z = expensive_calculation();  
    ... //  
}
```

As since C++11 compilers are required to wrap the necessary protection around the initialisation of static locals, also this is guaranteed to work:*

```
class Singleton {  
    ... //  
public:  
    static Singleton &getInstance() {  
        static Singleton instance;  
        return instance;  
    }  
};
```

*: Non-believers should consider to copy the above code, paste it to <https://gcc.godbolt.org/> (or similar) after adding a member with a runtime-dependant initialisation, and view the assembler output ...

Notifications with Condition Variables

A well-known abstraction in concurrent programming are combining mutexes with a signalling mechanism.

One main use of condition variables is to **avoid busy waiting** in producer-consumer designs,

- where consumer and producer run concurrently,
- exchanging data over some buffer data structure.



For more information on condition variables see:

http://en.cppreference.com/w/cpp/thread/condition_variable

Condition Variable Example (1)

The following *RingBuffer* class can put condition variables to good use ...

```
template<typename T, std::size_t N>
class RingBuffer {
    std::array<T, N+1> buf;
    std::size_t p = 0, g = 0;
    bool empty() const { return p == g; }
    bool full() const { return (p+1) % buf.size() == g; }
public:
    void put(const T &val) {
        if (full())
            ... // handle case no space is available
        buf[p++] = val; p %= buf.size();
    }
    void get(T &val) {
        if (empty())
            ... // handle case no data is available
        val = buf[g++]; g %= buf.size();
    }
};
```

... exactly at the currently omitted points.

Condition Variable Example (2)

Obviously there are two conditions, that need special attention:

- The buffer may be full when put is called, or
- it may be empty, when get is called.

Therefore two condition variables* are added, furthermore a mutex to protect accessing the buffer:

```
class RingBuffer {
    ... //
    ... // as before
    ... //
    std::condition_variable data_available;
    std::condition_variable space_available;
    std::mutex buffer_access;
public:
    ... // see next page
};
```

*: As the buffer space cannot be full and empty at the same time, technically one condition variable would suffice, but for this introductory example using two different instances seems to be clearer.

Condition Variable Example (3)

There are two operations (of interest here), applicable to condition variables, **sending** and **waiting for** notifications:^{*}

```
class RingBuffer
    ... // see previous page
public:
    void put(const T &val) {
        std::unique_lock<std::mutex> lock(buffer_access);
        space_available.wait(lock, [this]{ return !full(); });
        buf[p++] = val; p %= buf.size();
        data_available.notify_one();
    }
    void get(T &val) {
        std::unique_lock<std::mutex> lock(buffer_access);
        data_available.wait(lock, [this]{ return !empty(); });
        val = buf[g++]; g %= buf.size();
        space_available.notify_one();
    }
};
```

^{*}: Essential here is also the connection between the condition variables, the mutex protecting the RingBuffer invariants, and the conditions checked as part of waiting, which are detailed on the next page.

Waiting Anatomy

Waiting on a condition variable – as shown on the last page – with

```
// as part of put:
std::unique_lock<std::mutex> lock(buffer_access);
space_available.wait(lock, [this]{ return !full(); });

// as part of get:
std::unique_lock<std::mutex> lock(buffer_access);
data_available.wait(lock, [this]{ return !empty(); });
```

is equivalent to the following, **with the mutex being locked before:**

```
// as part of put:
while (full()) { // !!full()
    buffer_access.unlock();
    // wait for notification
    buffer_access.lock();
}

// as part of get:
while (empty()) { // !!empty
    buffer_access.unlock();
    // wait for notification
    buffer_access.lock();
}
```

At this point **the mutex is locked** (again) and **the condition is true.**

Spurious Wakeups

In the example code before, the loop (in the equivalent of wait-ing) may seem unnecessary, as the respective notification will be sent only after some action has made the condition true.

Nevertheless it makes sense and may even be necessary:

- **First of all, an implementation is allowed to give **Spurious Wake-Ups**, so the loop is necessary anyway.**
- If notifications are sent while nobody waits on the condition variable, it is simply discarded, therefore
 - a producer-consumer scenario is more robust if it tends to send "too many" notifications (of which some are discarded) ...
 - ... while sending "too few" could cause some thread to wait forever.

Specifying the condition check in combination with wait-ing *does it right* and hence should be preferred over writing a loop explicitly.*

*: Thus avoiding some later maintainer considers the while "unnecessary" and replace it with if ...

Atomic Operation Support

With the support for atomic operations C++11 multi-threading allows to implement

- Wait-Free Algorithms,
- Lock-Free Algorithms, and
- Obstruction-Free Algorithms.

The basic concept is to provide a way to know whether a certain modification of some memory location was caused by the current thread or by another one.



For more information about atomic operation support see:
<http://en.cppreference.com/w/cpp/atomic>

Atomic Operations Example

The following example, demonstrating the lock-free approach with operations (instead of using mutexes) modifies a [former example](#):

```
template<typename T1, typename T2, typename Transformation>
void worker(T1 args[], std::size_t data_size, T2 result,
            std::atomic_size_t &total_progress,
            Transformation func) {
    while (total_progress < data_size) {
        constexpr auto chunks = std::size_t{100};
        std::size_t beg = total_progress.load();
        std::size_t end;
        do {
            end = ((data_size - beg) > chunks)
                ? beg + chunks
                : data_size;
        } while (!total_progress.compare_exchange_weak(beg, end));
        std::transform(&data[beg], &data[end], &sqrts[beg], func);
    }
}
```

Be sure to understand that the loop controlled by the return value of `compare_exchange_weak` guarantees the prior calculations are (still) valid.

Memory Order (and Dependencies)

Memory order "*specifies how regular, non-atomic memory accesses are to be ordered around an atomic operation*".*

There are a number of different models to select from,

- with the default providing "*sequentially consistent ordering*",
- being closest to what most developers would expect,
- but (possibly) not with the optimal performance for a given use case.



For more information on memory order see:

http://en.cppreference.com/w/cpp/atomic/memory_order

*: Cited from the reference further down on this page.

Atomic Operations Recommendation

- Except for the potential of deadlocks the challenges are similar to algorithms using locks:^{*}
 - Problems may only show for a critical timing and may be reproducible in particular test environments only.
 - In general, a failed test may show the presence of errors, but even many successful tests do not guarantee their absence.



Beyond trivial cases – like the one shown in the example –, implementing multi-threaded programs with atomic operations requires **substantial expertise**.

Be sure to keep the design **as simple as possible** and have it reviewed by other developers experienced in that particular fields, maybe both, colleagues and hired consultants too.

^{*}: It may very well be the case that problematic situations depend on hardware features like the size of cache-lines, the depth of an instruction pipeline, or the way branch prediction works.

Using Class `std::thread`

In the examples before the class `std::thread` was used only indirectly (via `std::async`, building on [Futures and Promises](#)).

- There is also a class `std::thread`
- taking any runnable code as constructor argument,
- executing it in a separate thread.



For more information on the `std::thread` class see:
<http://en.cppreference.com/w/cpp/thread>

There are explicitly no means provided to forcefully terminate one thread from another one.



Any non-portable way* to forcefully terminate a thread risks to entail serious consequences later, as e.g. locks may not be released or awaited notifications not be sent.

*: Such as potentially existing *interrupt* or *kill* functions reachable via a [Native Interface Handle](#).

Example for Using Class `std::thread`

In case worker **returns no value** and **throws no exception**, like in a former example (and assuming the same set-up), using class `std::thread` directly can be straight forward (left) or done in the "nifty" way (right):

```
// starting some workers
// multi-threaded ...
std::thread t1{worker_task};
std::thread t2{worker_task};
...
// ... and waiting for them
// to finish:
t1.join();
t2.join();
...

using namespace std;
constexpr auto NCORES = 4;
// starting one worker-thread
// per core ...
array<thread, NCORES> threads;
for (auto &t : threads)
    t = thread{worker_task};
// ... and wait for all to
// finish:
for (auto &t : threads)
    t.join();
```

An alternative to `join`-ing with a thread is `detach`-ing it.



A program will immediately terminate if an instance of class `std::thread` referring to an active thread gets destructed.

Recommendations for Using Class `std::thread`

- In trivial cases (no exceptions, no result to fetch) using threads via instances of class `std::thread` may be considered.
- Nevertheless understand the peculiarities and know how to avoid race conditions, especially when `std::thread` objects go out of scope.*



It causes the program to terminate if the callable started throws an exception or an instance of `std::thread` is still in *joinable state* when it goes out of scope and gets destructed.

*: At 1:00:47 the following video by Scott Meyers gives a good introduction to the problem and some recipes how it can be avoided: <http://channel9.msdn.com/Events/GoingNative/2013/An-Effective-Cpp11-14-Sampler>

Native Handles

Last and finally, most C++11 multi-threading implementation build on a threading model provided by their execution environment.

The requirements in the standard are rather the intersection of the features provided by well-known threading models.

Usually and typically the standard does not mandate any any extensions thereof, but in some cases provides a way to "reach through" to the native thread model:

- E.g. `std::thread::native_handle` could provide ways to manipulate thread priorities, maybe including ways to specify protocols for [Priority Ceiling](#) or other means to circumvent [Priority Inversion](#).
- Also the classes for condition variables and the various kinds of mutexes have member functions `native_handle`.
- The enumeration `std::launch` may provide more (named) values to control implementation specific details in the behavior of `std::async`.

Concurrency Recommendations

So far the presentation of C++11 concurrency support was only meant as an overview.



Practically using concurrency features beyond parallelizing independent tasks requires much more knowledge and experience in this area, what this presentation can not provide.

A good and near to exhaustive coverage of the concurrency part of C++11 is:

C++ Concurrency in Action
Practical Multithreading
by Anthony Williams
ISBN-13: 978-1-9334988-77-1

Boost: Threads Library

As C++11 threads emerged from [Boost.Thread](#) there is little difference.

Some differences between C++11 and Boost may exist (depending on the version of the latter) and new features may appear and be tried in Boost first before eventually getting part of an upcoming C++ standard.

There are many reasons why a pure library solution may have practical difficulties to properly support concurrency.

Most important is in the memory model, which needs to be clearly defined so that optimising compilers are restricted to the necessary limits – but not (too far) beyond these.*

*: A related, easy to recognize problem that demonstrates why concurrency cannot be added as a library alone but must be part of the core language part, can be seen when considering a long-existing features of C which is also part of C++: the initialization of local `static` variables. By definition this takes place when the definition is reached in the program flow for the first time and hence, in the general case must be protected with a mutex (or similar) to avoid race conditions (except for initialisation with a compile-time constant that can be loaded at program startup).